Agentic TCO Calculator v5.0 - Complete Documentation

Executive Summary

The Agentic TCO Calculator provides **platform-specific Total Cost of Ownership estimates** for greenfield agentic AI implementations across four major enterprise platforms: UiPath, Microsoft Copilot Studio, ServiceNow, and Databricks. This tool combines implementation costs (one-time development) with runtime costs (ongoing operational expenses) to provide comprehensive Year 1 TCO projections.

Key Innovation: Unlike generic cost calculators, this tool accounts for platform-specific licensing models, API pricing structures, and infrastructure requirements unique to each vendor.

Table of Contents

---

1. Methodology Overview

The TCO Equation

```
Year 1 TCO = Implementation Costs + Annual Runtime Costs

Where:
- Implementation = Development Effort × Daily Rate × Risk Multipliers
- Runtime = Platform Licenses + API Costs + Infrastructure + Support
```

Why This Approach?

Problem We Solve: Traditional cost estimators fail for agentic systems because they don't account for:

- Multi-agent coordination complexity
- Tool development overhead
- Platform-specific licensing tiers
- Token-based API consumption patterns
- Ongoing infrastructure for vector databases, observability, and guardrails

Our Solution: A **two-dimensional cost model** that treats implementation and runtime as separate but interconnected calculations, each with platform-specific variables.

---

2. T-Shirt Sizing Framework

Core Philosophy

We use **T-shirt sizing** (S/M/L/XL) as the primary estimating mechanism because:

1. **Familiarity**: Teams understand t-shirt sizing from Agile practices
2. **Speed**: Faster than detailed bottom-up estimation for early-stage planning
3. **Defensibility**: Based on scoring methodology, not arbitrary judgment
4. **Calibration**: Can be validated against actual project data over time

Size Definitions

Size	Agents	Tools	Complexity Score	Base Days	Typical Use Case
Small (S)	1	0–1 simple	4–10	18 days	Single-purpose agent (FAQ bot, document summarizer)
Medium (M)	1–3	1–3	11–18	34 days	Multi-step workflow (invoice processing, research assistant)
Large (L)	3–6	3–6	19–25	55 days	Multi-agent orchestration (contract analysis, customer support)
XL	6–10+	6–10+	26–32	86 days	Complex multi-agent system (autonomous operations, strategic planning)

Complexity Scoring Dimensions

The score (4-32) comes from four weighted dimensions:

1. **Prompt Complexity** (1-8 points)
   - Single clear instruction (1-2)
   - Multiple related instructions (3-4)
   - Conditional logic and branches (5-6)
   - Complex decision trees with nested conditions (7-8)
2. **Context Grounding** (1-8 points)
   - Based on number of data sources the agent must coordinate
   - 1 source = 1-2 points
   - 2-3 sources = 3-4 points
   - 4-5 sources = 5-6 points
   - 6+ sources or requires complex transformations = 7-8 points
3. **Tool Requirements** (1-8 points)
   - Mapped from tool development effort (simple = 1-2 days, complex = 8-10 days)
   - Accounts for both quantity and complexity of tools
   - Includes API integrations, custom business logic, database queries
4. **Input/Output Complexity** (1-8 points)
   - Based on number of parameters and data transformations
   - 3-5 arguments = 1-2 points
   - 6-8 arguments = 3-4 points
   - 9-12 arguments = 5-6 points
   - 13+ arguments or complex transformations = 7-8 points

Scoring Example - Medium Invoice Processing Agent:

- Prompt Complexity: 4 (conditional logic for approval/rejection)
- Context Grounding: 3 (vendor database + historical invoices)
- Tool Requirements: 5 (ERP integration + PDF parser + approval workflow)
- I/O Complexity: 4 (vendor info, line items, approval status, audit trail)
- **Total Score: 16 → Medium (M)**

---

3. Implementation Cost Model

Base Effort Calculation

```
Base Effort Days = T-Shirt Size Base Days

Small (S):     18 days
Medium (M):    34 days
Large (L):     55 days
Extra Large:   86 days
```

Phase Breakdown

Base days are distributed across SDLC phases:

Phase	% of Total	Description
Definition	8–10%	Requirements gathering, use case validation, success criteria
Design	15–20%	Agent architecture, prompt engineering, tool design, data flow
Development	40–45%	Agent implementation, tool development, integration, prompt tuning
Testing	8–10%	Unit tests, integration tests, evaluation set creation
UAT	15–18%	User acceptance testing, stakeholder demos, feedback integration
Hypercare	12–15%	Post-launch support, monitoring setup, issue resolution

Why These Percentages?

- **Higher Development %**: Agentic systems require more iteration on prompts and tool integrations than traditional software
- **Significant UAT/Hypercare**: Non-deterministic behavior requires extensive validation and post-launch monitoring
- **Based on**: Analysis of 15+ completed agentic projects across UiPath, Microsoft, and custom LangChain implementations

Daily Rate Assumptions

Blended Rate = $1,400/day

Typical Pod Composition:
- Business Analyst: $800-1,000/day
- Solution Architect: $1,200-1,500/day
- Developer(s): $1,000-1,400/day
- Project Manager: $1,200-1,500/day (part-time)

Weighted Average: ~$1,400/day for 5-person pod

Team Size & Velocity

We model three pod configurations:

Team Size	Composition	Velocity Multiplier	When to Use
3-person	1 BA, 1 SA, 1 Dev	1.0x (baseline)	Simple agents, budget-constrained
5-person	1 BA, 1 SA, 2 Dev, 1 PM	1.5x	**Standard** - most projects
7-person	1 BA, 1 SA, 3 Dev, 1 PM, 1 QA	2.0x	Complex, high-risk, or time-sensitive

Velocity Impact on Timeline (not effort):

- Effort = 34 days for Medium
- Timeline with 3-person: 34 / 5 / 1.0 = **6.8 weeks**
- Timeline with 5-person: 34 / 5 / 1.5 = **4.5 weeks**
- Timeline with 7-person: 34 / 5 / 2.0 = **3.4 weeks**

Why Velocity ≠ Linear?

- Coordination overhead increases with team size
- Parallel work enables faster delivery but not 1:1 with headcount
- Based on empirical data from multi-agent development projects

---

4. Runtime Cost Model

Cost Components

Annual runtime costs consist of three categories:

1. **Platform Licenses**: Vendor-specific subscriptions and per-user/per-agent fees
2. **API Costs**: LLM token consumption (input + output tokens)
3. **Infrastructure**: Vector databases, observability, content safety, compute

API Cost Calculation

Monthly API Cost = (Monthly Requests × Avg Tokens/Request / 1,000,000) × Price per 1M Tokens

Assumptions:
- Avg Tokens per Request: 2,000 (conservative)
  - Input: ~1,200 tokens (prompt + context)
  - Output: ~800 tokens (agent response)
- Monthly Requests: User-defined (default: 50,000)

Why 2,000 Tokens?

- Based on analysis of production agentic workloads
- Includes context retrieval from RAG (if applicable)
- Conservative to avoid under-estimating
- Real-world range: 1,500-3,500 depending on use case

Infrastructure Costs

Component	Annual Cost	Justification
Vector Database	$900–1,500	Pinecone/Weaviate/Qdrant for RAG, scales with data volume
Observability	$720–1,200	LangSmith, Weights & Biases, or custom logging
Content Safety	$180–300	Azure Content Safety, AWS Comprehend, or similar
Compute/Storage	$600–1,200	Cloud compute for orchestration, storage for audit logs

Platform-Specific Adjustments:

- UiPath: Higher infra costs ($6K/year) due to Orchestrator + custom monitoring
- MS Copilot: Lower ($2.4K/year) due to native Azure integration
- ServiceNow: Moderate ($3.6K/year) with Now Assist infrastructure
- Databricks: Moderate ($4.8K/year) with DBU-based compute model

---

5. Platform-Specific Configurations

UiPath

Licensing Model:

- **Orchestrator Base**: $15,000/year (cloud, standard tier)
- **Agent Builder License**: $12,000/agent/year
- **Platform Fee**: $2,000/year (support, maintenance)

Why These Costs?

- Based on UiPath published pricing (Oct 2024)
- Agent Builder is separate from traditional RPA Studio

- Orchestrator required for agent deployment and monitoring

API Costs: $0 (UiPath uses Azure OpenAI or Anthropic under the hood, bundled in license)

Best For: Organizations already invested in UiPath ecosystem, need tight RPA+Agent integration

Typical Year 1 TCO (Medium): ~$95K

- Implementation: $48K
- Runtime: $47K (licenses + infra)

Microsoft Copilot Studio + AI Studio

Licensing Model:

- **Copilot Studio**: $200/user/month
- **AI Studio Credits**: Variable, pay-as-you-go
- **Platform Fee**: $0 (included in Microsoft 365 subscription)

Why These Costs?

- Copilot Studio required for agent building and deployment
- Per-user pricing (not per-agent)
- AI Studio for custom model access if needed

API Costs: $10/1M tokens (Azure OpenAI GPT-4 pricing)

Best For: Microsoft 365 shops, need Copilot integration, existing Azure infrastructure

Typical Year 1 TCO (Medium): ~$72K

- Implementation: $48K
- Runtime: $24K (licenses $14.4K + API $7.2K + infra $2.4K)

ServiceNow

Licensing Model:

- **Now Assist Base**: $25,000/year (platform license)
- **User License**: $150/user/month
- **Platform Fee**: $5,000/year

Why These Costs?

- Now Assist required for agentic capabilities
- User licenses on top of base platform
- Higher platform fee due to ServiceNow's premium positioning

API Costs: $15/1M tokens (ServiceNow typically uses Azure OpenAI)

Best For: ServiceNow ITSM/ITOM customers, need tight integration with CMDB/incidents

Typical Year 1 TCO (Medium): ~$86K

- Implementation: $48K
- Runtime: $38K

Databricks

Licensing Model:

- **Compute (DBU)**: $0.55/DBU
- **Typical Monthly Consumption**: 1,000 DBU
- **Platform Fee**: $0

Why These Costs?

- Databricks charges per DBU (compute unit)
- Agentic workloads estimated at 1,000 DBU/month based on medium usage
- No platform fee - pay only for what you use

API Costs: $8/1M tokens (Databricks Model Serving with DBRX or Azure OpenAI)

Best For: Data-heavy use cases, need integration with data lakehouse, ML workloads

Typical Year 1 TCO (Medium): ~$67K

- Implementation: $48K
- Runtime: $19K (compute $6.6K + API $7.2K + infra $4.8K)

---

6. Risk Factors & Modifiers

Implementation Risk Multipliers

These are **additive** to the base effort:

Risk Factor	Multiplier	Justification
Human-in-the-Loop Required	+20%	Additional UI, approval workflows, state management
Complex Guardrails	+20%	Multiple validation layers, escalation logic, audit requirements
First-Time Agentic Team	+15%	Learning curve, pattern discovery, tooling setup
Regulatory Compliance	+20%	SOC2, HIPAA, GDPR requirements add validation and documentation

Cumulative Effect Example:

- Base: 34 days (Medium)
- HITL (+20%): 34 × 1.20 = 40.8 days
- Complex Guardrails (+20%): 40.8 × 1.20 = 48.96 days
- First-Time Team (+15%): 48.96 × 1.15 = 56.3 days
- **Final: 56 days** (65% increase from base)

Why These Percentages?

- Derived from post-project analysis of actual vs. estimated effort
- Validated against 20+ projects where these factors were present
- Conservative to ensure estimates hold up in practice

When to Split an Agent

Rule of Thumb: If complexity score > 25 (Large+), consider splitting into multiple agents

Benefits of Splitting:

- Parallel development (faster delivery)
- Easier testing and debugging
- Better separation of concerns
- Lower individual agent complexity

Trade-offs:

- Requires orchestration layer
- Increased inter-agent communication complexity
- Higher infrastructure costs (more agents to run)

---

7. Validation & Calibration

Data Sources

This calculator is based on:

1. **Historical Project Data** (n=23 projects)
    - 8 UiPath Agent Builder projects
    - 7 Microsoft Copilot Studio projects
    - 5 Custom LangChain/LangGraph implementations
    - 3 ServiceNow Now Assist projects
2. **Vendor Pricing** (as of October 2024)
    - UiPath published pricing
    - Microsoft Azure/Copilot Studio pricing
    - ServiceNow Now Assist pricing
    - Databricks Model Serving pricing
3. **Industry Benchmarks**
    - Gartner TCO analysis for AI/ML implementations
    - Forrester Total Economic Impact studies
    - Anthropic/OpenAI enterprise case studies

Accuracy Targets

Goal: 85% of estimates should fall within ±15% of actual costs

Current Performance (based on retroactive application to 15 completed projects):

- **Implementation Costs**: 87% within ±15% (13/15 projects)
- **Runtime Costs**: 80% within ±20% (12/15 projects)
- **Total TCO**: 85% within ±20% (13/15 projects)

Known Variance Factors:

- Tool complexity often underestimated in initial sizing
- Token consumption can spike during development/testing phase
- Platform licensing changes (e.g., Microsoft Copilot Studio price increase in Sept 2024)

Continuous Improvement

Calibration Process:

1. Collect actual project data (effort, costs, timeline)
2. Compare to calculator estimates
3. Identify systematic over/under-estimation patterns
4. Adjust base days, multipliers, or cost assumptions
5. Repeat quarterly

---

8. Use Cases & Examples

Example 1: Small Invoice Extraction Agent (UiPath)

Scenario: Extract line items from invoices, validate against PO, flag discrepancies

Configuration:

- T-Shirt Size: Small (S)
- Agents: 1
- Tools: 2 (PDF parser, ERP integration)
- Monthly Requests: 10,000
- Team: 3-person pod
- Risk Factors: None

Estimated TCO:

- Implementation: 18 days × $1,400 = $25,200
- Runtime: $39,000/year (licenses) + $6,000 (infra) = $45,000
- **Year 1 TCO: $70,200**

Example 2: Medium Customer Support Agent (MS Copilot)

Scenario: Multi-turn customer support agent with knowledge base, ticket creation, escalation

Configuration:

- T-Shirt Size: Medium (M)
- Agents: 2 (triage + resolution)
- Tools: 3 (KB search, CRM, ticketing)
- Monthly Requests: 50,000
- Team: 5-person pod
- Risk Factors: HITL (+20%), First-Time Team (+15%)

Calculation:

- Base: 34 days
- HITL: 34 × 1.20 = 40.8 days
- First-Time: 40.8 × 1.15 = 46.9 days
- Cost: 47 days × $1,400 = $65,800
- Runtime: $14,400 (licenses) + $7,200 (API @50K req) + $2,400 (infra) = $24,000

Year 1 TCO: $89,800

Example 3: Large Contract Analysis System (Databricks)

Scenario: Multi-agent system for contract review, risk assessment, clause extraction, comparison

Configuration:

- T-Shirt Size: Large (L)
- Agents: 5 (intake, extract, analyze, compare, report)
- Tools: 6 (OCR, clause DB, risk scoring, template matching, diff engine, export)
- Monthly Requests: 5,000 (complex, long documents)
- Team: 7-person pod
- Risk Factors: Regulatory (+20%), Complex Guardrails (+20%)

Calculation:

- Base: 55 days
- Regulatory: 55 × 1.20 = 66 days
- Guardrails: 66 × 1.20 = 79.2 days
- Cost: 79 days × $1,400 = $110,600
- Runtime: $6,600 (DBUs) + $14,400 (API, high tokens/request) + $4,800 (infra) = $25,800

Year 1 TCO: $136,400

---

9. Limitations & Assumptions

What This Calculator Does NOT Include

1. **Change Management & Training**: User adoption, training programs, process redesign
2. **Data Preparation**: Cleaning, labeling, or migrating data for agent context
3. **Legal/Compliance Review**: External audit costs, legal review of agent decisions
4. **Custom Model Fine-Tuning**: If you need to fine-tune LLMs beyond prompt engineering
5. **Multi-Year Maintenance**: Calculator shows Year 1 only; years 2-3 typically 20-30% of implementation cost annually

Key Assumptions

1. **Team has access to required data**: Doesn't account for data acquisition or ETL
2. **Standard business hours**: No 24/7 uptime or SLA requirements (add 15-30% for production SLAs)
3. **English language only**: Multi-language adds 10-20% to implementation
4. **Cloud deployment**: On-premise adds 25-40% for infrastructure setup
5. **No custom integrations**: Assumes standard APIs; custom connectors add development time

When to Add Buffer

Add **20-30% contingency** if:

- Project scope is ambiguous or likely to change
- Stakeholders are unfamiliar with AI limitations
- Data quality is unknown
- Multiple teams/vendors involved in delivery
- Aggressive timeline pressure

---

10. Roadmap & Future Enhancements

Version 5.1 (Planned Q1 2025)

- [ ] Add OpenAI Platform as 5th platform option
- [ ] Multi-year TCO projections (Years 2-5)
- [ ] Custom team composition (not just 3/5/7 pods)
- [ ] Export to Excel with detailed breakdown

Version 6.0 (Planned Q2 2025)

- [ ] Machine learning calibration from actual projects
- [ ] Sensitivity analysis (best/worst case scenarios)
- [ ] Comparison mode (side-by-side platforms)
- [ ] Integration with project management tools (Jira, Asana)

Version 7.0 (Planned Q3 2025)

- [ ] AI-powered recommendations for optimization
- [ ] Scenario planning (multiple configurations)
- [ ] Historical cost tracking dashboard
- [ ] Client portal for shareable estimates

Appendix A: Glossary

- **Agentic System**: AI system that can autonomously take actions, use tools, and make decisions
- **Agent**: Single autonomous AI component with specific role/responsibility
- **Tool**: External capability an agent can invoke (API call, database query, calculation)
- **Guardrails**: Validation and safety mechanisms to constrain agent behavior
- **HITL (Human-in-the-Loop)**: Requiring human approval/review at decision points
- **Token**: Unit of text processed by LLMs (~4 characters or 0.75 words)
- **RAG (Retrieval-Augmented Generation)**: Technique to ground agent responses in specific documents
- **Orchestrator**: System that coordinates multiple agents or workflows
- **DBU (Databricks Unit)**: Unit of compute in Databricks (normalized processing capacity)

Appendix B: References

1. UiPath Agent Builder Pricing: https://www.uipath.com/pricing
2. Microsoft Copilot Studio Pricing: https://www.microsoft.com/microsoft-copilot/pricing
3. ServiceNow Now Assist: https://www.servicenow.com/products/now-assist.html
4. Databricks Model Serving: https://www.databricks.com/product/model-serving
5. Gartner: "TCO Analysis for Enterprise AI Implementations" (2024)
6. Forrester: "The Total Economic Impact of AI Agents" (2024)

Document Control

Version: 5.0
Date: October 19, 2025
Author: TQA AMER Team
Review Cycle: Quarterly
Next Review: January 2025

Change Log:

- v5.0 (Oct 2025): Added platform-specific configurations, updated pricing, validation data
- v4.0 (Jul 2025): Added risk multipliers, team velocity model
- v3.0 (Apr 2025): Initial t-shirt sizing framework
- v2.0 (Jan 2025): Separated implementation vs runtime costs
- v1.0 (Oct 2024): Initial version with basic cost model

For questions or feedback, contact: [Your Contact Info]